

CA1

February 10, 2023

1 Compulsory Assignment 1 - Pandas and visualizations

1.0.1 Imports

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

1.1 Loading and exploring the dataset

1. Load the dataset named `airbnb.csv` and store it in a dataframe called `raw_df`. Use the column named `id` as the index column for the dataframe

```
[ ]: # Insert your code below
# =====

raw_df = pd.read_csv('assets/airbnb.csv', index_col='id', )
```

2. Print the first five rows of the dataframe

```
[ ]: # Insert your code below
# =====

print(raw_df.head(5))
```

	name	host_id	host_name	\
id				
183319	Panoramic Ocean View Venice Beach	867995	Barbara X	
109	Amazing bright elegant condo park front *UPGRA...	521	Paolo	
51307	Spanish Bungalow Guest House LA CA. 30 plus ni...	235568	David	
184314	Boho Chic Flat..Steps to Beach!	884031	Ashley	
51498	Guest House With Its Own Entrance/Exit and Hot...	236758	Bay	

	neighbourhood	latitude	longitude	room_type	price	\
id						
183319	Venice	33.99211	-118.47600	Entire home/apt	152	
109	Culver City	33.98301	-118.38607	Entire home/apt	115	

51307	Atwater Village	34.12206	-118.26783	Entire home/apt	75
184314	Venice	33.97487	-118.46312	Entire home/apt	125
51498	Mar Vista	34.00389	-118.44126	Entire home/apt	189

	minimum_nights	number_of_reviews	calculated_host_listings_count	\
id				
183319	30	3		2
109	30	2		1
51307	30	138		2
184314	30	30		1
51498	3	378		1

	availability_365	number_of_reviews_ltm	state	city
id				
183319	0	0	CA	Los Angeles
109	139	0	CA	Los Angeles
51307	224	0	CA	Los Angeles
184314	0	0	CA	Los Angeles
51498	348	41	CA	Los Angeles

3. How many unique values exist in each of the columns state and city?

```
[ ]: # Insert your code below
# =====

state = raw_df.state
city = raw_df.city

unique_state = state.nunique()
unique_city = city.nunique()

print(f'Number of unique states: {unique_state} \n')
print(f'Number of unique cities: {unique_city} \n')
```

Number of unique states: 19

Number of unique cities: 31

4. Identify missing (NaN) values in each of the columns in the dataset

```
[ ]: # Insert your code below
# =====

print(f'Column      Number of missing values ')
for c in raw_df.columns:
    n_NaN = raw_df[c].isnull().sum()
    print(f'{c:<32} {n_NaN}')
```

Column	Number of missing values
name	19

host_id	0
host_name	1144
neighbourhood	712
latitude	0
longitude	0
room_type	0
price	0
minimum_nights	0
number_of_reviews	0
calculated_host_listings_count	0
availability_365	0
number_of_reviews_ltm	0
state	0
city	0

5. Create a copy of `raw_df` named `df`. Remove any rows containing NaN values in the new dataframe. What is the shape of `df` before and after removing the NaN values?

```
[ ]: # Insert your code below
# =====

df = raw_df.copy()

df_shape_with_nan = df.shape

df.dropna(inplace=True)

print('The shape before removing NaN values: ', df_shape_with_nan)
print('The shape after removing NaN values: ', df.shape)
```

The shape before removing NaN values: (325858, 15)

The shape after removing NaN values: (323983, 15)

6. Which `room_type`, `state` and `city` is the most popular (by number of instances)? Print the name and count of each

Hint: The output should look something like this:

Column: [col], Most popular: [name], Count: [count]

Column: [col], Most popular: [name], Count: [count]

Column: [col], Most popular: [name], Count: [count]

```
[ ]: # Insert your code below
# =====

room_type = df.room_type
city = df.city
state = df.state

most_frequent_city = city.value_counts().index[0]
```

```

max_count_city = max(city.value_counts())

most_frequent_state = state.value_counts().index[0]
max_count_state = max(state.value_counts())

most_frequent_room_type = room_type.value_counts().index[0]
max_count_room_type = max(room_type.value_counts())

print(f'Column: {city.name:<11} Most popular: {most_frequent_city:<17} Count:␣
↪{max_count_city}')
print(f'Column: {state.name:<11} Most popular: {most_frequent_state:<17} Count:␣
↪{max_count_state}')
print(f'Column: {room_type.name:<11} Most popular: {most_frequent_room_type:
↪<17} Count: {max_count_room_type}')

```

```

Column: city           Most popular: Los Angeles           Count: 91600
Column: state          Most popular: CA                     Count: 127206
Column: room_type      Most popular: Entire home/apt       Count: 241433

```

7. What is the average and median price for a listing?

```

[ ]: # Insert your code below
# =====

price = df.price

print('The median of the price is: ', price.median())
print('The mean of the price is: ', price.mean())

```

```

The median of the price is: 159.0
The mean of the price is: 285.125163974653

```

8. What is the average price for the states CA, FL and NY?

Hint: The output should look something like this:

```

State: [col], Average price: [price]
State: [col], Average price: [price]
State: [col], Average price: [price]

```

```

[ ]: # Insert your code below
# =====

CA = df[df.state == 'CA']
mean_price_CA = CA.price.mean()

FL = df[df.state == 'FL']
mean_price_FL = FL.price.mean()

NY = df[df.state == 'NY']

```

```
mean_price_NY = NY.price.mean()

print(f'State: CA      Average price: {mean_price_CA}')
print(f'State: FL      Average price: {mean_price_FL}')
print(f'State: NY      Average price: {mean_price_NY}')
```

```
State: CA      Average price: 288.39531940317283
State: FL      Average price: 241.98664420647336
State: NY      Average price: 197.21922246220302
```

9. Create a new dataframe called `df_beach` containing all listings with “beach” in the name. Print out the shape of `beach_df`

The filtering should not be case sensitive, meaning that names containing beach, Beach, BeAcH etc. all should be included

```
[ ]: # Insert your code below
# =====

df_beach = df[df['name'].str.contains('beach', case=False)]

print(df_beach.shape)
```

```
(31436, 15)
```

1.2 Visualizing the dataset

10. Create plot with 2 vertical axes and one horizontal axes. The plot should display a barchart containing the count of the 10 most popular states and cities, each in its own subplot. The bars should be sorted in descending order.

Use `df` in all tasks in this section

Hint: It is recommended to use the `Barplot` function built into Seaborn for barcharts.

The output should look something like this:


PS: Disregard the color scheme of the example image.

```
[ ]: # Insert your code below
# =====

states = df.state.value_counts()[:10]
cities = df.city.value_counts()[:10]

fig, axs = plt.subplots(2)

sns.barplot(ax=axs[0], x=states.index, y=states)
sns.barplot(ax=axs[1], x=cities.index, y=cities)
axs[1].tick_params(axis='x', rotation=45)
```



11. Create a scatterplot with the longitude and latitude of the listings in df. Longitude should be on the x-axis and latitude on the y-axis.

The output should look something like this:

PS: Disregard the color scheme of the example image.


```
[ ]: # Insert your code below
# =====

longitude = df.longitude
latitude = df.latitude

plt.scatter(longitude, latitude)

plt.xlabel('longitude')
plt.ylabel('latitude')
```

```
[ ]: Text(0, 0.5, 'latitude')
```




12. Create a matrix containing the correlations between the different columns in `df`. Plot it as a heatmap using Seaborn or similar. What does the plot tell you about correlations? Which columns are the most correlated to price?

```
[ ]: # Insert your code below
# =====

corr = df.corr(numeric_only = True)
sns.heatmap(corr)
```

```
[ ]: <AxesSubplot: >
```



The plot is a heatmap, where the lighter the pixel is the higher the correlation is.

Price is best correlated with calculated_host_listing_count and slightly correlated with availability_365.