

DAT200 CA5 2022

Kaggle username: jorid holmen

Imports

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.impute import SimpleImputer
import csv
from sklearn.metrics import r2_score

from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.decomposition import PCA
from sklearn.linear_model import LinearRegression
```

Reading data

```
In [ ]: train = pd.read_csv('data/train.csv', index_col=0)
test = pd.read_csv('data/test.csv', index_col=0)
```

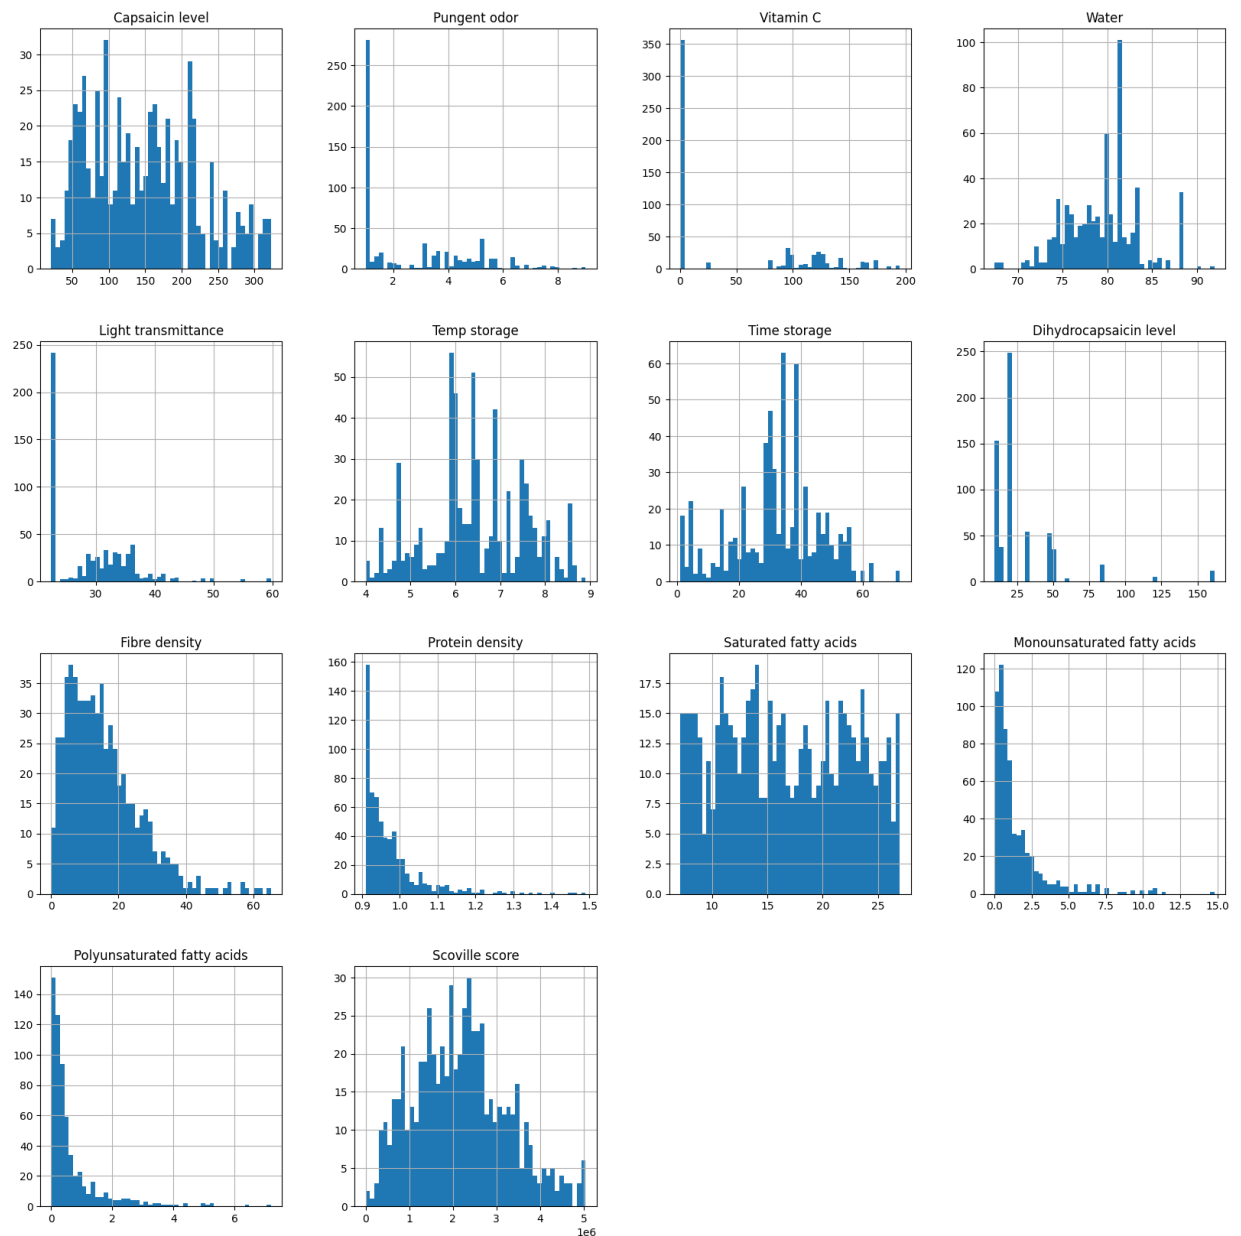
Data exploration and visualisation

```
In [ ]: train.head()
```


```
Out[ ]:
```

	Capsaicin level	Pungent odor	Vitamin C	Water	Light transmittance	Temp storage	Time storage	Dihydrocapsaicin level
0	166.7	6.8	0.0	77.95	32.4	7.5	4.0	32.
1	170.4	5.7	0.0	74.06	39.0	4.7	52.0	47.
2	212.5	3.1	0.0	88.27	22.4	5.9	1.0	161.
3	216.0	1.0	0.0	80.01	22.4	6.4	31.0	10.
4	146.5	1.3	141.0	78.64	28.8	4.6	38.0	20.

```
In [ ]: train.hist(figsize=(20,20), bins=50)
plt.show()
```



```
In [ ]: sns.heatmap(train.corr())
plt.show()
```



Data cleaning

```
In [ ]: # checking for NaN values
print(f'Column Number of missing values in train data ')
for c in train.columns:
    n_NaN = train[c].isnull().sum()
    print(f'{c:<32} {n_NaN}')

print(f'Column Number of missing values in test data')
for c in test.columns:
    n_NaN = test[c].isnull().sum()
    print(f'{c:<32} {n_NaN}')
```

```
Column Number of missing values in train data
Capsaicin level          0
Pungent odor             0
Vitamin C                0
Water                   0
Light transmittance      0
Temp storage             0
Time storage             0
Dihydrocapsaicin level  0
Fibre density            0
Protein density          0
Saturated fatty acids    0
Monounsaturated fatty acids 0
Polyunsaturated fatty acids 0
Scoville score           0
Column Number of missing values in test data
Capsaicin level          0
Pungent odor             0
Vitamin C                0
Water                   0
Light transmittance      0
Temp storage             0
Time storage             0
Dihydrocapsaicin level  0
Fibre density            0
Protein density          0
Saturated fatty acids    0
Monounsaturated fatty acids 0
Polyunsaturated fatty acids 0
```

```
In [ ]: # checking for outliers
        # fjerne kategorier som har lav korrelasjon med scoville
```

Data exploration after cleaning

```
In [ ]:
```

Data preprocessing

```
In [ ]:
```

Train test split

```
In [ ]: # splitting into X and y
        y = train['Scoville score']
        X = train.iloc[:,0:13]

        # Split data into training and test data
        X_train, X_test, y_train, y_test = train_test_split(
            X, y, test_size=0.3, random_state=1)
```

Scaling

```
In [ ]: # scaling in pipeline
```

Modelling

Data pipeline with regression model

```
In [ ]: pipe_rfreg = make_pipeline(RandomForestRegressor())

param_grid = {'randomforestregressor__random_state': list(np.arange(100,1000)),
              'randomforestregressor__n_estimators': list(np.arange(1,10))

gs_lr = GridSearchCV(estimator=pipe_rfreg,
                     param_grid=param_grid,
                     scoring='neg_mean_squared_error',
                     cv=10,
                     n_jobs=-1)

gs_lr_test = gs_lr.fit(X_train, y_train)
clf_rfr_best = gs_lr_test.best_estimator_
clf_rfr_best.fit(X_train, y_train)
print(r2_score(y_test, clf_rfr_best.predict(X_test)))
print(gs_lr_test.best_estimator_)

0.7921992956541402
Pipeline(steps=[('randomforestregressor',
                  RandomForestRegressor(n_estimators=9, random_state=300))
                ]))
```

```
In [ ]: r = RandomForestRegressor(n_estimators=9, random_state=300)
r.fit(X_train, y_train)
print(r2_score(y_test, r.predict(X_test)))

0.7921992956541402
```

```
In [ ]: clf_rfr_best.fit(X, y)
y_pred = clf_rfr_best.predict(test)

# write the results to a csv file
with open('kaggle_submission_rfr.csv', 'w') as f:
    w = csv.writer(f)

    w.writerow(['Id', 'Scoville score'])

    for r in range(0, 412):
        w.writerow([r, int(y_pred[r])])
```

Data pipeline with classification model

```
In [ ]: y_train_binned = pd.cut(y_train, 10, labels=False)
y_test_binned = pd.cut(y_test, 10, labels=False)

pipe_ada = make_pipeline(StandardScaler(),
                        AdaBoostClassifier())

param_grid = {'adaboostclassifier__estimator': [DecisionTreeClassifier()],
              'adaboostclassifier__n_estimators': list(np.arange(1,10)),
              'adaboostclassifier__learning_rate': [0.05, 0.1, 0.15, 0.2],
              'adaboostclassifier__random_state': list(np.arange(100,1001))}

gs_ada = GridSearchCV(estimator=pipe_ada,
                      param_grid=param_grid,
                      scoring='neg_mean_squared_error',
                      cv=10,
                      n_jobs=-1)

gs_ada_test = gs_ada.fit(X_train, y_train_binned)
clf_ada_best = gs_ada_test.best_estimator_
clf_ada_best.fit(X_train, y_train_binned)
print(clf_ada_best.score(X_test, y_test_binned))
print(gs_ada_test.best_estimator_)

0.3602150537634409
Pipeline(steps=[('standardscaler', StandardScaler()),
                 ('adaboostclassifier',
                  AdaBoostClassifier(estimator=DecisionTreeClassifier(),
                                     learning_rate=0.05, n_estimators=1,
                                     random_state=800))])
```

```
In [ ]: clf_ada_best.fit(X, y)
y_pred = clf_ada_best.predict(test)

# write the results to a csv file
with open('kaggle_submission_ada.csv', 'w') as f:
    w = csv.writer(f)

    w.writerow(['Id', 'Scoville score'])

    for r in range(0, 412):
        w.writerow([r, int(y_pred[r])])
```

Other models used for Kaggle submission

```
In [ ]:
```

Final Evaluation

```
In [ ]:
```

Kaggle submission

```
In [ ]: r = RandomForestRegressor(n_estimators=9, random_state=300)
r.fit(X, y)
y_pred = r.predict(test)

# write the results to a csv file
with open('kaggle_submission.csv', 'w') as f:
    w = csv.writer(f)

    w.writerow(['Id', 'Scoville score'])

    for r in range(0, 412):
        w.writerow([r, int(y_pred[r])])
```